

University of Washington Bothell
CSS503: System Programming
Program 1
Parallelizing a Convex-Hull Program

1. Purpose

In this programming assignment, we will parallelize a convex hull program with multiple processes. The algorithm that we use is a combination of divide-and-conquer and Graham's scan. Divide-and-conquer divides an entire problem into a tree of sub-spaces, each small enough to be solved with an independent process. We fork child processes in a tree as partitioning a problem with divide-and-conquer and let these processes work on their own sub-space in parallel.




Figure 1.

2. Convex Hull

Definition: The convex hull $CH(Q)$ of a set Q of points is the smallest convex polygon P for which each point in Q is either on the boundary of P or in its interior. Intuitively, we can think of each point in Q as being a nail sticking out from a board. The convex hull is then the shape formed by a tight rubber band that surrounds all the nails (from Introduction to Algorithms by T. H. Cormen et al).




Figure 2. The Convex Hull solution for a given set of points

3. Graham's Scan

This algorithm maintains a stack S of candidate points. Each point of the input set Q is pushed once onto the stack, and the points that are not vertices of $CH(Q)$ are eventually popped from the stack. When the algorithm terminates, stack S contains exactly the vertices of $CH(Q)$, in counterclockwise order of their appearance on the boundary. The scan will be performed as follows:

- Choose three consecutive points, last2, last1, and next;
- check if the leftward scanning bar that emanates from last2 actually hits last1 first before the next point;
- if so keeps last1 as a candidate, otherwise drop off last1;
- shift last1 to next (i.e., old next will be last1);
- and get a new point from the stack S as next.
- repeat until we reach the very first point.




Figure 3. Example of Graham's Scan

This link will be very useful for understanding Graham's Scan:

<https://www.geeksforgeeks.org/convex-hull-set-2-graham-scan/>

4. Combination with Divide-and-Conquer

Graham's Scan completes in $O(n \log n)$ but runs sequentially. We will parallelize this program using divide-and-conquer that recursively divides Q into two subsets until each becomes small enough to have at least two or more points and thus to create a triangle or a line segment quickly as a convex hull. Thereafter, we will go through the conquer phase as repeatedly merging two convex hulls into a larger subset, excluding inner points within this union of the hulls, and applying Graham's Scan to it at each repetitive stage.




Figure 4. Dividing the Problem in Figure 2

5. Program Structure

Your work will be to parallelize the convex hull program provided. The code to start from has been posted on Canvas.

The convex.cpp file consists of the following functions:

Function Name	Description
void init (deque<Point> &q, int nPoints)	Initializes points, each with (x, y) axes and an id.
void divide (deque<Point> &q, deque<Point> &s1, deque<Point> &s2)	Divides q into s1 and s2. It is used in divide_and_conquer().
double polar_angle (Point &o, Point &p, double &distance)	Computes the polar angle and distance between the origin o and the other point p.
void merge (deque<Point> &q, deque<Point> &s1, deque<Point> &s2)	Merges two deques s1 and s2 into q as removing those points not in polar angles.
bool leftturn (Point &a, Point &b, Point &c)	Checks if B is hit by line that revolves left from point A to point C.
void graham (deque<Point> &q)	Performs the graham algorithm that scans all points from the bottom as revolving a line leftward and eliminating non-convex points.
void divide_and_conquer (deque<Point> &q)	Divides a list of all points by two until each subset includes at least two, and thereafter merges those that make a convex.
int main (int argc, char *argv[])	Creates and initialize a given number (argv[1]) of points, starts a timer, calls divide_and_conquer() to find the convex hull for these points, and stops a timer to check the execution time.

6. Parallelization with Multi-Processes

Let's assume that you would like to parallelize this convex-hull program, (i.e., convex.cpp) with **N** Processes. When you recursively divide Q into two smaller subsets, (say left and right subsets), you should fork a child process and let it work on the right subset, whereas have the original parent take care of the left subset. If you still need to create a child process, have these parent and child processes create a new child respectively. This process **fork** should be performed within the `divide_and_conquer()` function of the original program, convex.cpp. Since all point data including global and local variables are automatically copied to new children, you don't have to worry about any data transfer from a parent to new children. Therefore, you don't even have to modify any functions other than `divide_and_conquer()` of the original convex.cpp.

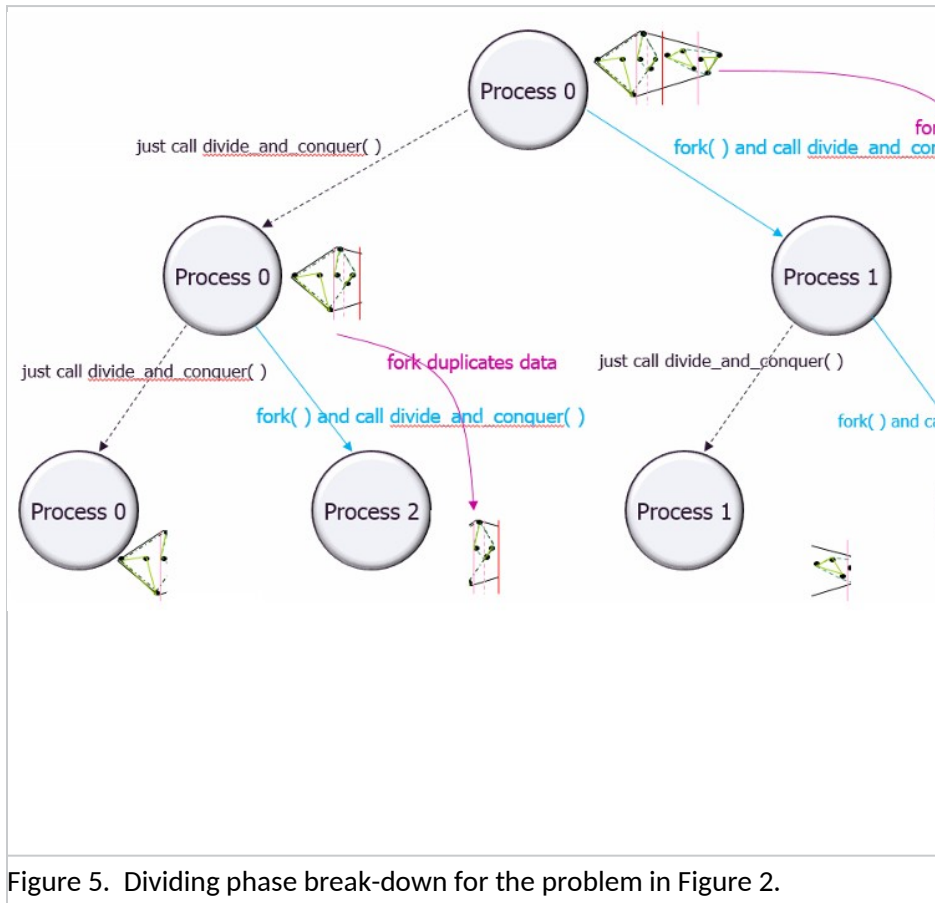


Figure 5. Dividing phase break-down for the problem in Figure 2.

However, in the conquer phase where a parent receives a sub convex hull from a child, you need to use a pipe that allows the child to send its sub convex hull back to its parent. For this purpose, every time a parent forks a new child, you should create a pipe in advance so that the parent and child process can share the same pipe. This pipe creation and data transfer from the child back to the parent should be also implemented within `divide_and_conquer()`.




Figure 6. Conquering Phase for the problem in Figure 2, continuing from the dividing phase in Figure 5 above.

When sending a sub convex hull from a child to its parent, add the following `send()` and `recv()` functions to your program so that you can only focus on process management in `divide_and_conquer()`.

```
/*
 * Sends a deque to a destination process through a pipe
 * fd.
 * @param fd  the file descriptor of a
 * pipe
 * @param q    a deque to
 * send
 */
```

```
void send(int fd, deque<Point> &q)
{
    int size = q.size();
    double x[size];
    double y[size];
    int id[size];

    // serialize all deque items to x, y, and id
    arrays
    for ( int i = 0; i < size; i++ )
    {
        Point p = q.front( );
        x[i] = p.x;
        y[i] = p.y;
        id[i] = p.id;
        q.pop_front( );
    }
    // send all data through a
    pipe
    write(fd, &size, sizeof(int));
    write(fd, x, sizeof(double) * size);
    write(fd, y, sizeof(double) * size);
    write(fd, id, sizeof(int) * size);
}
```

```
/*
 * Sends a deque from a source process through a pipe
 * fd.
 * @param fd  the file descriptor of a
 * pipe
 * @param q    a deque to
 * receive
 */
```

```
void recv( int fd, deque<Point> &q )
{
    // receive all data through a
    pipe
    int size = 0;
    read(fd, &size, sizeof(int));

    double x[size];
    double y[size];
    int id[size];
    read(fd, x, sizeof(double) * size);
    read(fd, y, sizeof(double) * size);
}
```

```

read(fd, id, sizeof(int) * size);

// de-serialized x, y, and id arrays to all deque items
for ( int i = 0; i < size; i++ )
{
    Point *p = new Point(x[i], y[i], id[i]);
    q.push_back(*p);
}
}

```

7. Statement of Work

Follow the steps described below:

Step 0: Get convex.cpp to work and understand what it is doing.

Step 1: Copy the original convex.cpp into a new file convex_mp.cpp (this is the one you will parallelize)

Step 2: Add the send() and recv() functions just above divide_and_conquer() to convex_mp.cpp.

Step 3: Add the following header files to the top of convex_mp.cpp.

```

#include <sys/types.h> // fork, wait
#include <sys/wait.h>  // wait
#include <unistd.h>    // fork, pipe
#include <stdlib.h>     // exit
#include <stdio.h>      // perror

```

Step 4: Parallelize convex_mp.cpp by modifying the divide_and_conquer() as described above.

Step 5: Conduct a performance evaluation and write up your report.

Focus on **testing** using #process = 1, 2, and 4. This is indicated through the second argument for your program.

Run your program with the following scenarios and attach your execution output to your report. There are sample outputs below.

```

./convex 100
./convex_mp 100 4
./convex 20000
./convex 20000_mp 1
./convex 20000_mp 2
./convex 20000_mp 4

```



```

css503@uw1-320-18 prog1]$ ./convex 100
do you want to display initial data? n
elapsed time = 1032
do you want to display result data? y
point[47].x = 7739.000000 .y = 12.000000
point[72].x = 9503.000000 .y = 19.000000
point[27].x = 9956.000000 .y = 1873.000000
point[45].x = 9932.000000 .y = 5060.000000
point[74].x = 9708.000000 .y = 6715.000000
point[28].x = 6862.000000 .y = 9170.000000
point[99].x = 5928.000000 .y = 9529.000000
point[15].x = 3929.000000 .y = 9802.000000
point[64].x = 709.000000 .y = 8927.000000
point[32].x = 336.000000 .y = 6505.000000
point[85].x = 124.000000 .y = 4914.000000
point[53].x = 97.000000 .y = 2902.000000
point[33].x = 846.000000 .y = 1729.000000
point[50].x = 795.000000 .y = 570.000000
point[5].x = 2362.000000 .y = 27.000000
[css503@uw1-320-18 prog1]$ ./convex_mp 100 4
do you want to display initial data? n
elapsed time = 2386
do you want to display result data? y
point[47].x = 7739.000000 .y = 12.000000
point[72].x = 9503.000000 .y = 19.000000
point[27].x = 9956.000000 .y = 1873.000000
point[45].x = 9932.000000 .y = 5060.000000
point[74].x = 9708.000000 .y = 6715.000000
point[28].x = 6862.000000 .y = 9170.000000
point[99].x = 5928.000000 .y = 9529.000000
point[15].x = 3929.000000 .y = 9802.000000
point[64].x = 709.000000 .y = 8927.000000
point[32].x = 336.000000 .y = 6505.000000
point[85].x = 124.000000 .y = 4914.000000
point[53].x = 97.000000 .y = 2902.000000
point[33].x = 846.000000 .y = 1729.000000
point[50].x = 795.000000 .y = 570.000000
point[5].x = 2362.000000 .y = 27.000000
[css503@uw1-320-18 prog1]$ ./convex 20000
do you want to display initial data? n
elapsed time = 200391
do you want to display result data? n
[css503@uw1-320-18 prog1]$ ./convex_mp 20000 1
do you want to display initial data? n
elapsed time = 211026
do you want to display result data? n

```

Another set of results:

```
css503@uw1-320-03:~/programming/progl$ ./convex 20000 1
do you want to display initial data? n
elapsed time = 58602
do you want to display result data? n
css503@uw1-320-03:~/programming/progl$ ./convex_mp 20000 1
do you want to display initial data? n
elapsed time = 57407
do you want to display result data? n
css503@uw1-320-03:~/programming/progl$ ./convex_mp 20000 2
do you want to display initial data? n
elapsed time = 30208
do you want to display result data? n
css503@uw1-320-03:~/programming/progl$ ./convex_mp 20000 4
do you want to display initial data? n
elapsed time = 25871
do you want to display result data? n
```

Your minimum requirements to complete this assignment include:

1. Your program creates `argv[2]` - 1 child processes and involve all `argv[2]` processes in parallel computation.
2. The performance improvement with two processes applied to 20,000 points should be at least 1.7 times
3. The performance improvement with four processes applied to 20,000 points should be at least 2.3 times.

8. What to Turn in / Rubric

	Criteria	Points
Code	Source code that adheres good modularization, coding style, and an appropriate amount of comments. Follow the google C++ coding guidelines: https://google.github.io/styleguide/cppguide.html	5
	Execution output that verifies the correctness of your implementation and demonstrates any improvement of your program's execution performance.	5
	Correctness code builds and executes correctly as well as handles error conditions	10
Report	Design document of your parallelization strategies including explanations and illustration in one or two pages. (No more than two)	5
	Discuss Results <u>in one or two pages</u> . (No more than two) <ul style="list-style-type: none">• Analysis of the effectiveness of your parallelization.• Possible performance improvement of your program• Limitations of your program.	5
	Total	30

References

A couple of might-be-useful links discussing Convex Hull

<http://www.geeksforgeeks.org/convex-hull-set-2-graham-scan/>

<http://axon.cs.byu.edu/~martinez/classes/312/Projects/Project2/project2.html>